Community Detection Methods for the Generalized Random Dot Product Graph Model

A dissertation proposal submitted in partial satisfaction of the requirements for the degree of
Doctor of Philosophy
in
Statistical Science

John Koo

Research Committee Members

Dr. Michael Trosset
Dr. Minh Tang
Dr. Julia Fukuyama
Dr. Roni Khardon
Dr. Fangzheng Xie

June 2, 2021

Department of Statistics
Indiana University
Bloomington, Indiana

Contents

Abstract

Graph and network data, in which samples are represented not as a collection of feature vectors but as relationships between pairs of observations, are increasingly widespread in various fields ranging from sociology to computer vision. One common goal of analyzing graph data is community detection or graph clustering, in which the graph is partitioned into disconnected subgraphs in an unsupervised yet meaningful manner (e.g., by optimizing an objective function or recovering unobserved labels). Because traditional clustering techniques were developed for data that can be represented as vectors, they cannot be applied directly to graphs. In this research, we investigate the use of a family of spectral decomposition based approaches for community detection in block models (random graph models with inherent community structure), first by demonstrating how under the Generalized Random Dot Product Graph framework, all graphs generated by block models can be represented as feature vectors, then applying clustering methods for these feature vector representations, and finally deriving the asymptotic properties of these methods.

1 Introduction

1.1 Research Goal

Graph and network data have become increasingly widespread in various fields including sociology, neuroscience, biostatistics, and computer science. This has resulted in various challenges for researchers who rely on traditional statistical and machine learning methods, many of which are incompatible with graph data and instead require the data to exist as feature vectors in Euclidean space. This includes graph clustering and community detection, which are often the goal of data analysis on graphs and networks. Common clustering methods typically involve calculating some central or representative point for each cluster around which the data belonging to that cluster lie (e.g., Lloyd's algorithm for K-means clustering [5], Gaussian Mixture Models [3]). Because these methods involve computing summary statistics within each cluster, such as the sample average, they cannot be applied directly to graphs, necessitating methods for transforming the graph data into feature vectors if they are to be used.

One family of methods for unifying graph community detection with traditional clustering techniques is Spectral Clustering [12], which involves embedding the graph into Euclidean space, followed by applying a popular clustering algorithm such as K-means clustering. The Random Dot Product Graph (RDPG) [14] and Generalized Random Dot Product Graph (GRDPG) [9] models take this further by explicitly constructing generative models for graphs via latent vectors in Euclidean space. In such a model, each vertex has a corresponding vector in latent space, and the edge probabilities between pairs of vertices are computed by operations on the corresponding pairs of latent vectors. A community detection algorithm motivated by this may involve learning the latent positions given an observed graph and then learning the community labels given the latent positions.

The aim of our research is to develop consistent community detection techniques under the RDPG and GRDPG frameworks. First, we explore existing generative graph models with underlying community structures (Block Models) that can be inferred by connecting these models to the RDPG or GRDPG. In particular, we use the connection between the Popularity Adjusted Block Model [10] and the GRDPG to motivate two community detection methods. Then we explore other latent structures or mixture distributions in the latent space that induce graphs for which consistent community detection is possible.

1.2 Notation

Let $G = (V, E)$ be an undirected, unweighted graph with no self-loops and n vertices. Denote $A \in \{0,1\}^{n \times n}$ as the adjacency matrix of G such that $A_{ij} = 1$ if there exists an edge between vertices i and j and $A_{ij} = 0$ otherwise. Because G is symmetric and contains no self-loops, $A_{ij} = A_{ji}$ and $A_{ii} = 0$ for $i, j \in [n]$. We further restrict our analyses to independent Bernoulli graphs. Let $P \in [0,1]^{n \times n}$ be a symmetric matrix of edge probabilities. Graph G is sampled from P by drawing $A_{ij} \overset{\text{ind}}{\sim} \text{Bernoulli}(P_{ij})$ for each $1 \leq i < j \leq n$ ($A_{ji} = A_{ij}$ and $A_{ii} = 0$). We denote $A \sim \text{BernoulliGraph}(P)$ as graph with adjacency matrix A sampled from edge probability matrix P in this manner. If each vertex has a hidden label in $[K]$, they are denoted as $z_1, ..., z_n$. Finally, we denote $X = \begin{bmatrix} x_1 & \cdots & x_n \end{bmatrix}^\top \in \mathbb{R}^{n \times d}$ as the sample $x_1, ..., x_n \in \mathbb{R}^d$.

2 Preliminaries

2.1 Block Models

Generative models for independent Bernoulli graphs involve defining the edge probability matrix P whose ij^{th} entry is the probability of an edge between vertices i and j for each $i, j \in [n]$. In order to motivate community detection methods, we restrict the generative model such that for each pair of vertices (i, j), the probability of an edge between the vertices P_{ij} depends in some way on their labels z_i and z_j. Such models are called Block Models.

One type of Block Model is the Stochastic Block Model (SBM) [6]: Given K communities and each vertex belonging to one of the K communities, the SBM restricts P to $K(K+1)/2$ unique entries such that $P_{ij} = p_{z_i z_j}$ where $p_{kl} = p_{lk}$ is the probability of an edge between each vertex in community k to each vertex in community l. The homogeneous SBM further restricts P to two unique entries: $P_{ij} = p$ if $z_i = z_j$ and $P_{ij} = q$ otherwise. Several generalizations of the SBM have been introduced since, including the Degree Corrected Block Model (DCBM) [4] and the Popularity Adjusted Block Model (PABM) [10]. Like the SBM, these models involve imposing some restriction on P based on community labels (and possibly other factors).

2.2 (Generalized) Random Dot Product Graphs

Like the SBM, DCBM, and PABM, the RDPG and GRDPG are generative models for graphs involving an edge probability matrix P. The RDPG starts with points in latent space $X \in \mathcal{X} \subset \mathbb{R}^d$ for $\mathcal{X} = \{x, y \in \mathbb{R}^d : x^\top y \in [0, 1]\}$. P is then constructed as $P = XX^\top$, and $A \sim \text{BernoulliGraph}(P)$. We provide a more formal definition of the RDPG and GRDPG as follows:

Definition 1 ((Generalized) Random Dot Product Graph). *Let $X \in \mathbb{R}^{n \times d}$ be a collection of n points in $\mathcal{X} \subset \mathbb{R}^d$ where $\mathcal{X} = \{x, y \in \mathbb{R}^d : x^\top y \in [0, 1]\}$. $G = (V, E)$ is a Random Dot Product Graph if its adjacency matrix is drawn as $A \sim BernoulliGraph(XX^\top)$.*

If on the other hand, $\mathcal{X} = \{x, y \in \mathbb{R}^{p+q} : x^\top I_{p,q} y \in [0, 1]\}$ and $A \sim BernoulliGraph(XI_{p,q}X^\top)$ for $I_{p,q} = blockdiag(I_p, -I_q)$ and $p + q = d$, then A is the adjacency matrix of a Generalized Random Dot Product Graph. These are denoted by $A \sim RDPG(X)$ and $A \sim GRDPG_{p,q}(X)$ respectively.

In addition, let F be a probability distribution with support $\mathcal{X}$, and $X_1, ..., X_n \stackrel{iid}{\sim} F$ with $X = \begin{bmatrix} X_1 & \cdots & X_n \end{bmatrix}^\top$. If A is drawn from X as before, then $(A, X) \sim RDPG(F, n)$ or $(A, X) \sim GRDPG_{p,q}(F, n)$.

The structure of the RDPG and GRDPG provides a straightforward method for recovery of the latent positions via spectral decomposition.

Definition 2 (Adjacency Spectral Embedding). *Let $A \sim RDPG(X)$ for $X \in \mathcal{X} \subset \mathbb{R}^{n \times d}$. Let $A \approx V \Lambda V^\top$ be the approximate spectral decomposition of A corresponding to the d largest eigenvalues and their corresponding eigenvectors. Then the rows of $V \Lambda^{1/2}$ are the scaled Adjacency Spectral Embedding (ASE) of A, and the rows of V are the unscaled ASE of A.*

If $A \sim GRDPG_{p,q}(X)$, then let $A \approx V \Lambda V^\top$ be the approximate spectral decomposition of A corresponding to the p most positive and q most negative eigenvalues of A and their corresponding eigenvectors. Then the rows of $V |\Lambda|^{1/2}$ and V are the scaled and unscaled ASE of A respectively.

Athreya et al. [1] showed that under mild conditions, if $(A_n, X_n) \sim \text{RDPG}(F, n)$ and $\hat{X}_n$ is the scaled ASE of A_n, for some sequence of orthogonal matrices $\{W_n\}$,

$$\max_i \|(\hat{X}_n)_i - W_n(X_n)_i\| \overset{a.s.}{\to} 0 \tag{1}$$

Similarly, Rubin-Delanchy et al. [9] showed that for $(A_n, X_n) \sim \mathrm{GRDPG}_{p,q}(F, n)$,

$$\max_i \|(\hat{X}_n)_i - Q_n(X_n)_i\| \overset{a.s.}{\to} 0 \tag{2}$$

where $\{Q_n\}$ is a sequence of matrices in $\mathbb{O}(p, q)$, the indefinite orthogonal group of order p, q.

It is straightforward to show that all Bernoulli graphs with positive semidefinite P are special cases of the RDPG, which is a special case of the GRDPG, and all Bernoulli Graphs generated by P regardless of positive semidefiniteness are special cases of the GRDPG. This includes the SBM, DCBM, and PABM. In the following example, we show that under the RDPG framework, the latent configuration that induces the SBM takes on a very particular form.

Example (Connecting the SBM to the RDPG). *Let $G = (V, E)$ with adjacency matrix A be sampled from the SBM with two communities for which the within-edge probabilities are p and q for communities 1 and 2 respectively, and the between-community edge probability is r, with $pq > r^2$. Let community 1 have n_1 vertices, community 2 have n_2 vertices, and $n_1 + n_2 = n$. Without loss of generality, organize P and A such that the kl^{th} block represents edges between communities k and l. Then $P = \begin{bmatrix} P^{(11)} & P^{(12)} \\ P^{(21)} & P^{(22)} \end{bmatrix}$ where each block is a constant value, e.g., $P_{ij}^{(11)} = p$. Then one RDPG representation of this SBM is:*

$$X = \begin{bmatrix} \sqrt{p} & 0 \\ \vdots & \vdots \\ \sqrt{p} & 0 \\ \sqrt{r^2/p} & \sqrt{q - r^2/p} \\ \vdots & \vdots \\ \sqrt{r^2/p} & \sqrt{q - r^2/p} \end{bmatrix} \in \mathbb{R}^{n \times 2}$$

where the first n_1 rows are $\begin{bmatrix} \sqrt{p} & 0 \end{bmatrix}$ and the next n_2 rows are $\begin{bmatrix} \sqrt{r^2/p} & \sqrt{q - r^2/p} \end{bmatrix}$. This is shown by reconstructing P as $P = XX^\top$. Thus the SBM is equivalent to a RDPG with the latent configuration of two point masses, one for each community.

The ASE of the assortative SBM consists of points in $\mathbb{R}^K$ that lie near one of K centers, depending on the community label, leading to ASE followed by K-means clustering (or similar, e.g., GMM) as a consistent community detection algorithm [7]. A similar result has been shown for the DCBM [7], and our work extends this to the PABM.

2.3 Manifold Learning

Trosset et al. [11] showed that the ASE of a RDPG can be used to recover one-dimensional manifolds. Suppose $f : [0, 1] \mapsto \mathcal{X}$ such that f is smooth and $\mathcal{X}$ represents a curve or one-dimensional manifold in $\mathbb{R}^d$. If $t_1, ..., t_n \overset{\mathrm{iid}}{\sim} F$ such that F has support $[0, 1]$, the latent positions are $x_i = f(t_i)$ with y_i is its corresponding point in the scaled ASE, and $d_\epsilon(\cdot, \cdot)$ is the shortest path distance of an ϵ-neighborhood

graph. Under certain mild conditions, the shortest path distances of the ϵ-neighborhood graph of the ASE approaches the arc lengths along f:

$$d_\epsilon(y_i, y_j) \xrightarrow{p} \int_{t_i}^{t_j} \sqrt{\sum_r^d \left(\frac{df_r}{dt}\right)^2} dt \tag{3}$$

Athreya et al. [2] extended this further by generating a RDPG from a mixture of distributions on a curve. In their example, points were sampled from a mixture of two Beta distributions on the Hardy-Weinberg curve to construct the latent positions of a RDPG, with the goal of recovering the hidden mixture distribution from an observed graph.

3   The Popularity Adjusted Block Model

As discussed in §2.2, all Bernoulli graph models can be expressed as a RDPG or GRDPG model, and in particular, the (associative) SBM is a RDPG with a very specific structure in the latent space. A similar result has been shown for the DCBM [7]. We can also extend this to the PABM. In §3.1, we show that the PABM is equivalent to the GRDPG model for which the communities correspond to subspaces in the latent configuration. We further show that while the latent configuration is not unique, there exists one in particular for which the subspaces are orthogonal. This leads to two straightforward methods for community detection, one involving computing inner products in the ASE and another which is a straightforward application of an existing algorithm for subspace clustering. Finally, we will briefly explore alternative, non-spectral-based methods for community detection in the ASE involving Expectation Maximization and Bayesian methods.

3.1   Connecting the PABM to the GRDPG

We first define the PABM using the construction provided by Noroozi et al. [8].

Definition 3 (Popularity Adjusted Block Model). *Let $P \in [0, 1]^{n \times n}$ be a symmetric edge probability matrix for a set of n vertices, V. Each vertex has a community label $1, ..., K$, and the rows and columns of P are arranged by community label such that $n_k \times n_l$ block $P^{(kl)}$ describes the edge probabilities between vertices in communities k and l ($P^{(lk)} = (P^{(kl)})^\top$). Let graph $G = (V, E)$ be an undirected, unweighted graph such that its corresponding adjacency matrix $A \in \{0, 1\}^{n \times n}$ is a realization of BernoulliGraph(P).*

If each block $P^{(kl)}$ can be written as the outer product of two vectors:

$$P^{(kl)} = \lambda^{(kl)}(\lambda^{(lk)})^\top \tag{4}$$

for a set of K^2 fixed vectors $\{\lambda^{(st)}\}_{s,t=1}^K$ where each $\lambda^{(st)}$ is a column vector of dimension n_s, then graph G and its corresponding adjacency matrix A is a realization of a popularity adjusted block model with parameters $\{\lambda^{(st)}\}_{s,t=1}^K$.

We will use the notation $A \sim \text{PABM}(\{\lambda^{(kl)}\}_K)$ to denote a random adjacency matrix A drawn from a PABM with parameters $\lambda^{(kl)}$ consisting of K underlying communities.

It is straightforward to show that the PABM (as well as all Bernoulli Graphs) is a special case of the GRDPG. It can also be shown that the latent positions of the PABM under the GRDPG framework

4

consists of K K-dimensional subspaces in $\mathbb{R}^{K^2}$. While there is no unique latent configuration X such that $XX^\top = P$, the edge probability P for the PABM, they all have this subspace structure, and one in particular consists of *orthogonal* subspaces.

Theorem 1 (Connecting the PABM to the GRDPG for $K = 2$). *Let*

$$X = \begin{bmatrix} \lambda^{(11)} & \lambda^{(12)} & 0 & 0 \\ 0 & 0 & \lambda^{(21)} & \lambda^{(22)} \end{bmatrix}$$

$$U = \begin{bmatrix} 1 & 0 & 0 & 0 \\ 0 & 0 & 1/\sqrt{2} & 1/\sqrt{2} \\ 0 & 0 & 1/\sqrt{2} & -1/\sqrt{2} \\ 0 & 1 & 0 & 0 \end{bmatrix}$$

where each $\lambda^{(kl)}$ is a vector as in Definition 1. Then $A \sim GRDPG_{3,1}(XU)$ and $A \sim PABM(\{(\lambda^{(kl)}\}_2)$ are equivalent.

Theorem 2 (Generalization to $K > 2$). *There exists a block diagonal matrix $X \in \mathbb{R}^{n \times K^2}$ defined by PABM parameters $\{\lambda^{(kl)}\}_K$ and orthonormal matrix $U \in \mathbb{R}^{K^2 \times K^2}$ that is fixed for each K such that $A \sim GRDPG_{K(K+1)/2, K(K-1)/2}(XU)$ and $A \sim PABM(\{(\lambda^{(kl)}\})_K)$ are equivalent.*

Proof. Define the following matrices from $\{\lambda^{(kl)}\}_K$:

$$\Lambda^{(k)} = \begin{bmatrix} \lambda^{(k,1)} & \cdots & \lambda^{(k,K)} \end{bmatrix} \in \mathbb{R}^{n_k \times K}$$

$$X = \text{blockdiag}(\Lambda^{(1)}, ..., \Lambda^{(K)}) \in \mathbb{R}^{n \times K^2} \tag{5}$$

$$L^{(k)} = \text{blockdiag}(\lambda^{(1k)}, ..., \lambda^{(Kk)}) \in \mathbb{R}^{n \times K}$$

$$Y = \begin{bmatrix} L^{(1)} & \cdots & L^{(K)} \end{bmatrix} \in \mathbb{R}^{n \times K^2}$$

Then $P = XY^\top$.

Note that $Y = X\Pi$ for a permutation matrix Π, resulting in $P = X\Pi X^\top$. The permutation described by Π has K fixed points, which correspond to K eigenvalues equal to 1 with corresponding eigenvectors e_k where $k = r(K + 1) + 1$ for $r = 0, ..., K - 1$. It also has $\binom{K}{2} = K(K-1)/2$ cycles of order 2. Each cycle corresponds to a pair of eigenvalues $+1$ and -1 and a pair of eigenvectors $(e_s + e_t)/\sqrt{2}$ and $(e_s - e_t)/\sqrt{2}$.

Then Π has $K(K+1)/2$ eigenvalues equal to 1 and $K(K-1)/2$ eigenvalues equal to -1. Π has the decomposed form

$$\Pi = U I_{K(K+1)/2, K(K-1)/2} U^\top \tag{6}$$

The edge probability matrix then can be written as:

$$P = XUI_{p,q}(XU)^\top \tag{7}$$

$$p = K(K+1)/2 \tag{8}$$

$$q = K(K-1)/2 \tag{9}$$

and we can describe the PABM with K communities as a GRDPG with latent positions XU with signature $(K(K+1)/2, K(K-1)/2)$. $\qquad\square$

Example $(K = 3)$. *Using the same notation as in Theorem 2:*

$$X = \begin{bmatrix} \lambda^{(11)} & \lambda^{(12)} & \lambda^{(13)} & 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & \lambda^{(21)} & \lambda^{(22)} & \lambda^{(23)} & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 & \lambda^{(31)} & \lambda^{(32)} & \lambda^{(33)} \end{bmatrix}$$

$$Y = \begin{bmatrix} \lambda^{(11)} & 0 & 0 & \lambda^{(12)} & 0 & 0 & \lambda^{(13)} & 0 & 0 \\ 0 & \lambda^{(21)} & 0 & 0 & \lambda^{(22)} & 0 & 0 & \lambda^{(23)} & 0 \\ 0 & 0 & \lambda^{(31)} & 0 & 0 & \lambda^{(32)} & 0 & 0 & \lambda^{(33)} \end{bmatrix}$$

Then $P = XY^\top$ and $Y = X\Pi$ where Π is a permutation matrix consisting of 3 fixed points and 3 cycles of order 2:

$$\Pi = \begin{bmatrix} 1 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 1 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 & 1 & 0 & 0 \\ 0 & 1 & 0 & 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 1 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 & 0 & 1 & 0 \\ 0 & 0 & 1 & 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 1 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 1 \end{bmatrix}$$

Positions 1, 5, 9 are fixed.

The cycles of order 2 are $(2, 4)$, $(3, 7)$, and $(6, 8)$.

Therefore, we can decompose $\Pi = UI_{6,3}U^\top$ where the first three columns of U consist of e_1, e_5, and e_9 corresponding to the fixed positions 1, 5, and 9, the next three columns consist of eigenvectors $(e_k + e_l)/\sqrt{2}$, and the last three columns consist of eigenvectors $(e_k - e_l)/\sqrt{2}$, where pairs (k, l) correspond to the cycles of order 2 described above.

The latent positions are the rows of

$$XU = \begin{bmatrix} \lambda^{(11)} & 0 & 0 & \lambda^{(12)}/\sqrt{2} & \lambda^{(13)}/\sqrt{2} & 0 & \lambda^{(12)}/\sqrt{2} & \lambda^{(13)}/\sqrt{2} & 0 \\ 0 & \lambda^{(22)} & 0 & \lambda^{(21)}/\sqrt{2} & 0 & \lambda^{(23)}/\sqrt{2} & -\lambda^{(21)}/\sqrt{2} & 0 & \lambda^{(23)}/\sqrt{2} \\ 0 & 0 & \lambda^{(33)} & 0 & \lambda^{(31)}/\sqrt{2} & \lambda^{(32)}/\sqrt{2} & 0 & -\lambda^{(31)}/\sqrt{2} & -\lambda^{(32)}/\sqrt{2} \end{bmatrix}$$

Algorithm 1: Orthogonal Spectral Clustering.

Data: Adjacency matrix A, number of communities K

Result: Community assignments $1, ..., K$

1 Compute the eigenvectors of A that correspond to the $K(K+1)/2$ most positive eigenvalues and $K(K-1)/2$ most negative eigenvalues. Construct V using these eigenvectors as its columns.

2 Compute $B = |nVV^\top|$, applying $|\cdot|$ entry-wise.

3 Construct graph G using B as its similarity matrix.

4 Partition G into K disconnected subgraphs (e.g., using edge thresholding or spectral clustering).

5 Map each partition to the community labels $1, ..., K$.

In Theorem 2, we showed that a possible latent configuration for a K-community PABM under the GRDPG framework consists of vectors in $\mathbb{R}^{K^2}$ such that each community corresponds to a K-dimensional subspace and each subspace is orthogonal to the others. However, because latent configurations are not unique, we cannot apply this fact directly for community detection.

Remark. Let $A \sim \mathrm{GRDPG}_{p,q}(Y)$. Then $A \sim \mathrm{GRDPG}_{p,q}(YQ) \; \forall Q \in \mathbb{O}(p,q)$. This is because $YQI_{p,q}Q^\top Y^\top = YI_{p,q}Y^\top$.

However, we can show that if we treat the rows of the eigenvectors of P (ignoring the eigenvalues) as an embedding (unscaled ASE), then they always form orthogonal subspaces. This is stated more formally in Theorem 3, and we call the resulting algorithm Orthogonal Spectral Clustering.

Theorem 3. *Let $P = V\Lambda V^\top$ be the spectral decomposition of the edge probability matrix of a PABM. Define $B = nVV^\top$. Then $B_{ij} = 0 \; \forall i, j$ in different communities.*

If $\hat{V}$ is the *unscaled* ASE of A, then results from Rubin-Delanchy et al. [9] imply $n\hat{V}\hat{V}^\top \overset{a.s.}{\to} nVV^\top$. Then for vertices i and j belonging to separate communities, the ij^{th} entry of $n\hat{V}\hat{V}^\top$ approaches 0 with probability 1:

Theorem 4. *Let $\hat{B}_n$ with entries $\hat{B}_n^{(ij)}$ be the affinity matrix from OSC (Alg. 1). Then $\forall$ pairs (i, j) belonging to different communities and sparsity factor satisfying $n\rho_n = \omega\{(\log n)^{4c}\}$,*

$$\max_{i,j} |n(\hat{v}_n^{(i)})^\top \hat{v}_n^{(j)}| = O_P\left(\frac{(\log n)^c}{\sqrt{n\rho_n}}\right) \tag{10}$$

This provides the result that $\forall i, j$ in different communities, $\hat{B}_n^{(ij)} \overset{a.s.}{\to} 0$.

Since the communities are subspaces in the latent space, we can also use existing methods for subspace clustering on the ASE. Of particular interest is Sparse Subspace Clustering (SSC), which is performed by solving an optimization problem for each observed point in a sample. Given $X \in \mathbb{R}^{n \times d}$ with vectors $x_i^\top \in \mathbb{R}^d$ as rows of X, the optimization problem $c_i = \arg\min_c \|c\|_1$ subject to $x_i = X^\top c$ and $c^{(i)} = 0$ is solved for each $i \in [n]$. The solutions are collected into matrix $C = \begin{bmatrix} c_1 & \cdots & c_n \end{bmatrix}^\top$ to construct affinity matrix $B = |C| + |C^\top|$. If each x_i lie perfectly on one of K subspaces, B is sparse such that $B_{ij} = 0 \; \forall x_i, x_j$ belonging to different subspaces. Then B can describe a graph with at least K disjoint subgraphs, and if the number of subgraphs is exactly K, each subgraph maps onto a subspace.

In practice, SSC is performed by solving the LASSO problems:

$$c_i = \arg\min_c \frac{1}{2}\|x_i - X_{-i}^\top c\|_2^2 + \lambda\|c\|_1 \tag{11}$$

for some sparsity parameter $\lambda > 0$. The c_i vectors are then collected into C and B as described before. If X is noisy such that each x_i does not lie exactly on one of K subspaces but near it, the choice of λ becomes important in guaranteeing the Subspace Detection Property (SDP) [13].

Definition 4 (Subspace Detection Property). *Let $X = \begin{bmatrix} x_1 & \cdots & x_n \end{bmatrix}^\top$ be noisy points sampled from K subspaces. Let C and B be constructed from the solutions of LASSO problems as described in (11). If each column of C has nonzero norm and $B_{ij} = 0 \ \forall \ x_i$ and x_j sampled from different subspaces, then X obeys the Subspace Detection Property.*

Remark. In practice, a noisy sample X often does not obey SDP. In such cases, B is treated as an affinity matrix for a graph which is then partitioned into K subgraphs to obtain the clustering. On the other hand, if X does obey the SDP, B describes a graph with at least K disconnected subgraphs. Ideally, when SDP holds, there are exactly K subgraphs which map to each subspace, but it could be the case that some of the subspaces are represented by multiple disconnected subgraphs.

Since every ASE of the PABM consists of subspaces and as $n \to \infty$ each vector of the ASE approaches its subspace uniformly [9], SSC is also able to perform community detection for the PABM. It has been shown by Wang and Xu [13] that if the points lie sufficiently close to their respective subspaces and the cosine of the angles between subspaces is sufficiently small, SDP will hold. We now state that the unscaled ASE of the PABM exhibits exactly these conditions for sufficiently large n.

Theorem 5. *Let P_n describe the edge probability matrix of the PABM with n vertices, and let $A_n \sim Bernoulli(P_n)$. Let $\hat{V}_n$ be the matrix of eigenvectors of A_n corresponding to the $K(K+1)/2$ most positive and $K(K-1)/2$ most negative eigenvalues. Then $\exists \lambda > 0$ and $N < \infty$ such that when $n > N$, $\sqrt{n}\hat{V}_n$ obeys the subspace detection property with probability 1.*

3.2 Non-Spectral Approaches to Community Detection

So far, we connected the PABM to the GRDPG and used this to motivate community detection methods based on the ASE of the adjacency matrix sampled from the PABM. In this section, we will explore non-spectral methods for community detection.

Denoting $\{\lambda_{ik}\}$ as the scalar popularity parameters for the PABM such that $\lambda^{(kl)}$ is the vector of λ_{jl} terms for which vertex j is in community k, we can write the likelihood for the PABM:

$$p(A|z, \{\lambda_{ik}\}) = \prod_{i<j}(\lambda_{iz_j}\lambda_{jz_i})^{A_{ij}}(1 - \lambda_{iz_j}\lambda_{jz_i})^{1-A_{ij}}$$

$$= \prod_{i<j}\prod_k^K\prod_l^K(\lambda_{il}\lambda_{jk})^{A_{ij}z_{ik}z_{jl}}(1 - \lambda_{il}\lambda_{jk})^{(1-A_{ij})z_{ik}z_{jl}} \tag{12}$$

where $z_{ik} = I(Z_i = k)$.

The log-likelihood is

$$\log p = \sum_{i<j}\sum_k\sum_l z_{ik}z_{jl}(A_{ij}\log\lambda_{il}\lambda_{jk} + (1 - A_{ij})\log(1 - \lambda_{il}\lambda_{jk})) \tag{13}$$

8

3.2.1 Expectation Maximization

While direct maximization of the likelihood or log-likelihood is NP-complete, it is possible to apply an Expectation Maximization (EM) algorithm to this problem. To set this up, we write the complete data log likelihood as:

$$Q = \sum_{i<j} \sum_{k,l} z_{ik} z_{jl} (A_{ij} \log \lambda_{il}\lambda_{jk} + (1 - A_{ij}) \log(1 - \lambda_{il}\lambda_{jk})) + \sum_i \sum_k z_{ik} \log \pi_k$$

Where π_k are the *a priori* community label probabilities.

Then the expectation step involves solving for $\gamma_{ik} = P(Z_i = k \mid \{\pi_l\}, \{\lambda_{jl}\})$, which is given by

$$\log \gamma_{ik} \propto \log \pi_k + \sum_{j \neq i} \sum_l \pi_{jl} (A_{ij} \log \lambda_{il}\lambda_{jk} + (1 - A_{ij}) \log(1 - \lambda_{il}\lambda_{jk}))$$

For the maximization step, we can just set $\pi_k = \frac{1}{n} \sum_i \gamma_{ik}$, but the $\{\lambda_{ik}\}$ parameters cannot be maximized in closed form and require an optimization subroutine. Our proposed work involves finding an efficient optimization method for this subroutine.

3.2.2 Bayesian Methods

The introduction of community label probabilities $\pi_1, ..., \pi_K$ suggests a hierarchical model in which the community labels are generated by a prior distribution. We can apply the same principle to the popularity parameters $\{\lambda_{ik}\}$ to make the model fully Bayesian. In this setup, we sample $z_i \overset{\text{iid}}{\sim}$ Categorical$(\pi_1, ..., \pi_K)$ and $\lambda_{ik} \overset{\text{ind}}{\sim}$ Beta(a_{ik}, b_{ik}) before sampling the observed data $A_{ij} \overset{\text{ind}}{\sim} \lambda_{iz_j}\lambda_{jz_i}$.

The full joint distribution for this model can be written as:

$$\begin{aligned}
\log p = {}& constant \\
& + \sum_{i<j} \sum_k \sum_l z_{ik} z_{jl} (A_{ij} \log \lambda_{il}\lambda_{jk} + (1 - A_{ij}) \log(1 - \lambda_{il}\lambda_{jk})) \\
& + \sum_k \sum_i z_{ik} \log \pi_k \\
& + \sum_i \sum_k (a_{ik} - 1) \log \lambda_{ik} + (b_{ik} - 1) \log(1 - \lambda_{ik})
\end{aligned}$$

It is straightforward to obtain the conditional distribution for each z_i but not each λ_{ik}, so a Gibbs sampler cannot be derived for this model. However, we can use a Metropolis-Hastings algorithm. In particular, by taking a first-order Taylor approximation, we can approximate the conditional distribution for each λ_{ik} as a Beta distribution. Our proposed work involves investigating this method by simulating this model and implementing the Bayes sampler.

4   Generalizations to GRDPG-Based Community Detection

In the previous sections, we connected well-known and highly restricted generative graph models to the RDPG and GRDPG to show that highly structured latent configurations generate graphs

consistent with these models: The latent space for the SBM consists of K point masses, the latent space for the DCBM consists of K rays emanating from the origin, and the latent space for the PABM consists of K K-dimensional subspaces. In the following sections, we explore additional structured latent configurations corresponding to community structure and develop methods for community detection based on the consistency of the ASE and the structural forms of the latent configurations.

The general structure of interest can be described as follows: Suppose that in the latent space $\mathcal{X} \subset \mathbb{R}^d$, sample X of n points lie on a union of K disjoint manifolds with each manifold corresponding to a community. If $A \sim \mathrm{RDPG}(X)$, we wish to recover the community labels (up to permutation) from A.

Equivalently, suppose that probability distribution F is described as follows:

1. Define functions $f_1, ..., f_K$ such that $f_k : [0,1] \mapsto \mathcal{X}$ and $f_k(t) \neq f_l(t)$ $\forall k, l \in [K]$ and $t \in [0,1]$.
2. Sample labels $z_1, ..., z_n \overset{\text{iid}}{\sim} Categorical(\pi_1, ..., \pi_K)$.
3. Sample $t_1, ..., t_n \overset{\text{iid}}{\sim} D$ where D has support $[0,1]$.
4. Set latent positions $x_i = f_{z_i}(t_i)$ and $X = \begin{bmatrix} x_1 & \cdots & x_n \end{bmatrix}^\top$.

Then if $(A, X) \sim \mathrm{RDPG}(F, n)$ and we observe A, we wish to recover hidden labels $z_1, ..., z_n$.

4.1 Affine Segments

In this section, we will consider latent positions sampled uniformly from parallel unit length segments.

Example. *Let $U_1, ..., U_{n_1}, V_1, ..., V_{n_2} \overset{iid}{\sim} Uniform(0,1)$, $f_1(t) = (t, 0)$, and $f_2(t) = (t, a)$. $X_i = f_1(U_i)$ and $Y_j = f_2(V_j)$ $\forall i \in [n_1]$ and $j \in [n_2]$. If we observe $X_1, ..., X_{n_1}, Y_1, ..., Y_{n_2}$, what approach will allow us to group the observations by segment?*


Figure 1: Points sampled uniformly on two parallel unit length line segments

We will motivate an approach by the following example.

Example. *Let $U_1, ..., U_n \overset{iid}{\sim} Uniform(0,1)$ with order statistics $U_{(1)}, ..., U_{(n)}$. Then $\forall a \in (0,1)$ and $\delta \in (0,1)$, $\exists N = N(\delta, a) < \infty$ such that $\forall n \geq N$,*

$$P(\max_i U_{(i+1)} - U_{(i)} \leq a) \geq 1 - \delta/2 \tag{14}$$

Where $N(\delta, a)$ is monotone increasing with respect to δ and a. To prove this, we start with the fact that $U_{(i+1)} - U_{(i)} \sim Beta(1, n)$. Then

$$P(U_{(i+1)} - U_{(i)} \leq a) = 1 - (1-a)^n \tag{15}$$

and

$$P(\max_i U_{(i+1)} - U_{(i)} \leq a) \geq (P(U_{(i+1)} - U_{(i)} \leq a))^n = (1 - (1-a)^n)^n \quad (16)$$

This expression is monotone increasing $\forall n \geq N_1$ for some $N_1 < \infty$. Setting $(1-(1-a)^{N_2})^{N_2} \geq 1 - \delta/2$, we can solve for a finite N_2. Then $N = \max(N_1, N_2)$.

If we extend this example such that n_1 points are sampled uniformly from the segment $f_1(t) = (t, 0)$ and n_2 points are sampled uniformly from the segment $f_2(t) = (t, a)$ for $t \in [0, \cos\frac{\pi}{2}a]$, then a sample of size $\min(n_1, n_2) \geq N(\delta, a)$ is sufficient to satisfy:

$$P(\max_i X_{(i+1)} - X_{(i)} \leq a \leq \min_{i,j} \|X_i - Y_j\|) \geq 1 - \delta/2 \quad (17)$$

and similar for $P(\max_j Y_{(j+1)} - Y_{(j)} \leq \min_{i,j} \|X_i - Y_j\|)$, for X_i in the first segment and Y_j in the second segment and $X_{(i)}, Y_{(j)}$ are order statistics in the first coordinate. If each segment corresponds to a community, this leads to the following two results:

1. Single linkage clustering with $K = 2$ will perform perfect community detection with probability at least $1 - \delta$.

2. An ϵ-neighborhood graph with $\epsilon \in (0, a)$ will consists of at least 2 disjoint subgraphs such that no subgraph consists of members of two different communities (analogous to the SDP), with probability at least $1 - \delta$.

We can then further extend this to the case where points are drawn from unit segments with noise. It is straightforward to show that if the maximum noise is less than $a/3$, we can obtain a similar result as in the noiseless case. Thus if the distribution on the two segments is used to draw $(A, X) \sim \text{RDPG}(F, n)$, the same properties should hold for large n.

4.2 Manifolds

If instead of sampling uniformly from line segments of unit length, we sample uniformly from a 1 dimensional manifolds of unit length, the above property still holds. Let $U_1, ..., U_n \overset{iid}{\sim} Uniform(0, 1)$ and $f : [0, 1] \mapsto \mathbb{R}^d$ be a smooth function such that $\int_u^v \sqrt{\sum_k (df_k/dt)^2} dt = \|u - v\|$. Then $U_{(i+1)} - U_{(i)} \geq \|f(U_{(i+1)}) - f(U_{(i)})\|$, so $P(U_{(i+1)} - U_{(i)} \leq \alpha) \leq P(\|f(U_{(i+1)}) - f(U_{(i)})\| \leq a)$. If the shortest distance between the two manifolds defined by f_1 and f_2 with the same restriction is a, then the same N as before is sufficient, although perhaps a more lenient lower bound can be derived based on the shape of $f_k(\cdot)$.

Our proposed work extends these results by the following.

1. Show that the ASE of a (G)RDPG generated by this type of latent configuration produces the correct conditions as described above to apply the same probability statements.
2. Extend these results by deriving similar probability statements for nonuniform distributions.
3. Relax the condition for the minimum distance between manifolds by investigating alternative clustering techniques on the latent space (e.g., spectral clustering).

5 Summary

For this research, we propose using the connection between Block Models and the (Generalized) Random Dot Product Graph to motivate community detection methods. Previous research shows that the Stochastic Block Model and Degree Corrected Block Model have very structured latent configurations in the RDPG framework, and these latent structures can be used to develop community detection algorithms that result in perfect community detection for large n. In the same vein, we investigated the connection between the Popularity Adjusted Block Model and the GRDPG to motivate two community detection algorithms, the first of which is an application of an existing algorithm, Sparse Subspace Clustering, and the second is a new algorithm based on the particular latent form of the PABM, Orthogonal Spectral Clustering. Furthermore, we showed that for large n, our application of the SSC algorithm to the PABM upholds the Subspace Detection Property for large yet finite n, and the output of OSC results in perfect community detection almost surely as $n \to \infty$. We will extend these theoretical results with empirical simulations and real data examples. Finally, we will implement these two community detection algorithms in an R package.

We aim to extend these results by defining more general structures in the latent space to construct RDPG or GRDPG models and show that community detection is possible via Adjacency Spectral Embedding. So far, we showed that if in the latent space the communities can be represented as points sampled uniformly on curves or one-dimensional manifolds separated by a nonzero distance, we can apply two existing clustering techniques to the ASE and show that they will perform perfect community detection with high probability. We will generalize these results to nonuniform distributions on higher dimensional manifolds as well as investigate alternative clustering techniques for the latent space that may perform better for smaller n.

5.1 Estimated Timeline of Completion

Literature review: August 2021

Complete proofs of main theorems: January 2022

Simulations, real data analyses, R package: March 2022

Dissertation completion: April 2022

References

[1] Avanti Athreya, Donniell E. Fishkind, Minh Tang, Carey E. Priebe, Youngser Park, Joshua T. Vogelstein, Keith Levin, Vince Lyzinski, Yichen Qin, and Daniel L Sussman. Statistical inference on random dot product graphs: a survey. *Journal of Machine Learning Research*, 18(226):1–92, 2018. URL http://jmlr.org/papers/v18/17-448.html.

[2] Avanti Athreya, Minh Tang, Youngser Park, and Carey E. Priebe. On estimation and inference in latent structure random graphs, 2020.

[3] Chris Fraley and Adrian E Raftery. Model-based clustering, discriminant analysis, and density estimation. *Journal of the American Statistical Association*, 97(458):611–631, 2002. doi: 10.1198/016214502760047131. URL https://doi.org/10.1198/016214502760047131.

[4] Brian Karrer and M. E. J. Newman. Stochastic blockmodels and community structure in networks. *Physical Review E*, 83(1), Jan 2011. ISSN 1550-2376. doi: 10.1103/physreve.83.016107. URL http://dx.doi.org/10.1103/PhysRevE.83.016107.

[5] S. Lloyd. Least squares quantization in pcm. *IEEE Transactions on Information Theory*, 28 (2):129–137, 1982. doi: 10.1109/TIT.1982.1056489.

[6] François Lorrain and Harrison C. White. Structural equivalence of individuals in social networks. *The Journal of Mathematical Sociology*, 1(1):49–80, 1971. doi: 10.1080/0022250X.1971.9989788. URL https://doi.org/10.1080/0022250X.1971.9989788.

[7] Vince Lyzinski, Daniel L. Sussman, Minh Tang, Avanti Athreya, and Carey E. Priebe. Perfect clustering for stochastic blockmodel graphs via adjacency spectral embedding. *Electron. J. Statist.*, 8(2):2905–2922, 2014. doi: 10.1214/14-EJS978. URL https://doi.org/10.1214/14-EJS978.

[8] Majid Noroozi, Ramchandra Rimal, and Marianna Pensky. Estimation and clustering in popularity adjusted block model. *Journal of the Royal Statistical Society: Series B (Statistical Methodology)*, n/a(n/a). doi: https://doi.org/10.1111/rssb.12410. URL https://rss.onlinelibrary.wiley.com/doi/abs/10.1111/rssb.12410.

[9] Patrick Rubin-Delanchy, Joshua Cape, Minh Tang, and Carey E. Priebe. A statistical interpretation of spectral embedding: the generalised random dot product graph, 2017.

[10] Srijan Sengupta and Yuguo Chen. A block model for node popularity in networks with community structure. *Journal of the Royal Statistical Society. Series B: Statistical Methodology*, 80(2):365–386, March 2018. ISSN 1369-7412. doi: 10.1111/rssb.12245.

[11] Michael W. Trosset, Mingyue Gao, Minh Tang, and Carey E. Priebe. Learning 1-dimensional submanifolds for subsequent inference on random dot product graphs, 2020.

[12] Ulrike von Luxburg. A tutorial on spectral clustering. *CoRR*, abs/0711.0189, 2007. URL http://arxiv.org/abs/0711.0189.

[13] Yu-Xiang Wang and Huan Xu. Noisy sparse subspace clustering. *Journal of Machine Learning Research*, 17(12):1–41, 2016. URL http://jmlr.org/papers/v17/13-354.html.

[14] Stephen J. Young and Edward R. Scheinerman. Random dot product graph models for social networks. In Anthony Bonato and Fan R. K. Chung, editors, *Algorithms and Models for the Web-Graph*, pages 138–149, Berlin, Heidelberg, 2007. Springer Berlin Heidelberg. ISBN 978-3-540-77004-6.